

Huiqi Zou

Boston, MA | zou.huiqi@northeastern.edu | annazou1103.github.io

EDUCATION

Ph.D. in Computer Engineering

May 2030 (Expected)

Northeastern University, Advisor: Prof. Weiyang Shi

M.Sc.Eng in Computer Science

Dec. 2024

Johns Hopkins University, Thesis Advisor: Prof. Ziang Xiao

B.Sc. in Computer Science

June 2023

City University of Hong Kong, First Class Honors, Thesis Advisor: Prof. Linqi Song

RESEARCH INTERESTS

Evaluation and oversight of large language models (LLMs) and AI agents in human-centered settings.

- **Human oversight of AI agents:** detection and correction of unsafe or deceptive agent behavior.
- **Interactive evaluation:** multi-turn evaluation of LLMs, vision-language models, and chatbots.
- **User and persona simulation:** LLM-based simulation of human respondents and personalities.

PUBLICATIONS & PREPRINTS (*: equal authorship)

Publications

- Ye, J.*, **Zou, H.***, Yu, S., & Shi, W. (2026). Coding with “Enemy”: Can Human Developers Detect AI Agent Sabotage? *Advances in Neural Information Processing Systems (NeurIPS)* & **Best Paper Award at ICML DL4C**.
- Li, Y.*, **Zou, H.***, Wang, B., & Xiao, Z. (2026). Contextualized Evaluation of Vision Language Models through Dynamic, Multi-turn Interactions. *Advances in Neural Information Processing Systems (NeurIPS)* & *EMNLP GenBench*.
- Wang, P., **Zou, H.**, Jiang, H., Chen, H., Sun, T., Yi, X., Xiao, Z., & Oswald, F. L. (2026). Generative Personality Simulation via Theory-Informed Structured Interview. *Conference of the European Chapter of the Association for Computational Linguistics (EACL)*.
- **Zou, H.**, Wang, P., Yan, Z., Sun, T., & Xiao, Z. (2025). Can LLM “Self-report”? Evaluating the Validity of Self-report Scales in Measuring Personality Design in LLM-based Chatbots. *Conference on Language Modeling (COLM)*.
- Ma, X., **Zou, H.**, He, P., Keung, J., Li, Y., Yu, X., & Sarro, F. (2025). On the Influence of Data Resampling for Deep Learning-Based Log Anomaly Detection: Insights and Recommendations. *IEEE Transactions on Software Engineering*, 51(1), 243–261.
- Ma, X., Li, Y., Keung, J., Yu, X., **Zou, H.**, Yang, Z., Sarro, F., & Barr, E. T. (2025). Practitioners’ Expectations on Log Anomaly Detection. *IEEE Transactions on Software Engineering (TSE)*, 51(9), 2455–2471.
- Ma, X., Keung, J. W., Yu, X., **Zou, H.**, Zhang, J., & Li, Y. (2023). AttSum: A Deep Attention-based Summarization Model for Bug Report Title Generation. *IEEE Transactions on Reliability*, 72(4), 1663–1677.
- Liu, Z., Qin, Y., **Zou, H.**, Paek, E. J., Casenhiser, D., Zhou, W., & Zhao, X. (2022). Generating Natural Language Responses in Robot-Mediated Referential Communication Tasks to Simulate Theory of Mind. *International Conference on Social Robotics (ICSR)*, 100–109.

Preprints

- Ge, M., **Zou, H.**, Shi, W., & Jia, C. (2026). Conversations with AI Can Reduce Partisan Animosity by Correcting Misperceptions and Encouraging Personal Narratives. Under review.
- **Zou, H.**, Sun, Y., Keung, J., Li, Y., & Ma, X. (2026). Enriching Few-Shot Retrieval for Log Parsing via Structured Data Augmentation. Under review at *Information and Software Technology*.
- Wang, P., **Zou, H.**, Yan, Z., Guo, F., Sun, T., Xiao, Z., & Zhang, B. (2024). Not Yet: Large Language Models Cannot Replace Human Respondents for Psychometric Research. *OSF Preprints*.

RESEARCH EXPERIENCE

CHATS Lab, Northeastern University

Sep. 2025 – Present

Doctoral Researcher, Advisor: Prof. Weiyan Shi

- Designed a **realistic, long-horizon coding study** in which four AI models secretly attempted to sabotage developers' work during five-hour coding sessions with **over 100 developers** (NeurIPS & Best Paper Award @ ICML DL4C).
- Found that **94%** of developers failed to catch the sabotage, and **56%** accepted it after an AI safety monitor flagged it.

ISLE Lab, Johns Hopkins University

Oct. 2023 – May 2025

Research Assistant, Advisor: Prof. Ziang Xiao

- Built a **three-party framework for evaluating VLM hallucinations**: an examiner questions the model and a grader scores its answers over multi-turn conversations, uncovering more hallucinations than static methods (NeurIPS).
- Collected **500 human conversations** across five chatbot personality designs to compare **LLM personality self-ratings** with human-perceived ratings, finding limited alignment (COLM).

AiSE Research Group, City University of Hong Kong

July 2021 – May 2024

Undergraduate Researcher, Advisor: Prof. Jacky Keung

- Co-developed an **automated bug-report title generator** (RoBERTa encoder + Transformer decoder with a copy mechanism) that outperformed prior methods by **up to 58.8%** on text-quality metrics (IEEE Trans. Reliability).
- Implemented the **data-balancing pipeline** for a log anomaly detection study across four datasets, finding that **adding rare cases** is more effective than removing common ones (IEEE Trans. Software Engineering).

AI² Lab, City University of Hong Kong

Sep. 2022 – Apr. 2023

Undergraduate Researcher, Advisor: Prof. Linqi Song

- Trained a compact facial expression recognition model by **distilling knowledge** from a larger Swin Transformer V2 model, reducing computation while maintaining performance.

NICS, University of Tennessee

May 2022 – Aug. 2022

Undergraduate Researcher, Advisors: Prof. Xiaopeng Zhao and Dr. Kwai Lam Wong

- Collaborated with a team of two to integrate referential communication into feedback strategies, enhancing robot behaviors explainability and leveraging theory of mind to model human understanding.
- Developed a **BERT-based dialogue system** that extracts human perception from responses to generate contextually relevant feedback in referential communication tasks.

WORK EXPERIENCE

City University of Hong Kong, Research Assistant

Mar. 2023 - Jul. 2023

- Conducted sentiment and readability analysis on annual financial reports.

Siemens Limited (Hong Kong), Intern

June 2021 - Mar. 2022

- Developed the frontend of an Android app to engage audiences in science-related events.
- Maintained and updated an Angular website with a MongoDB-backed server for resource management.

TEACHING & SERVICE

- **Conference Reviewer:** EMNLP 2025, EACL 2026.
- **Workshop Area Chair:** IAB Workshop @ NeurIPS 2026.
- **Workshop Reviewer:** Multi-Turn Interactions in LLMs Workshop @ NeurIPS 2025, Meta-Agents Workshop @ NeurIPS 2026, UserSim @ NeurIPS 2026.
- **Course Assistant:** Intro to Human Language Technology, Johns Hopkins University, Fall 2024.
- **Student Mentor:** City University of Hong Kong, Fall 2021 – Spring 2023.

AWARDS AND HONORS

- **Best Paper Award**, ICML 2026 Workshop on Deep Learning for Code (DL4C) 2026
- **Second Prize of Huawei Final Year Project Competition**, City University of Hong Kong 2023
- **InfoTech Job Market Driven Scholarship**, InfoTech Services (Hong Kong) Limited 2023
- **Grand Prize Award & “Citi Challenge-ESG” Special Prize Award of 18th Citi Financial Innovation Application Competition**, Citi Bank (China) 2023
- **Reaching Out Award**, HKSAR Government Scholarship Fund 2022
- **Asia-Pacific Economic Cooperation Scholarship**, HKSAR Government Scholarship Fund 2022

TECHNICAL SKILLS

Proficient: Python, PyTorch, Hugging Face Transformers, TensorFlow, SLURM GPU clusters, Git

Prior experience: C++, Java, Kotlin, Docker, JavaScript/TypeScript